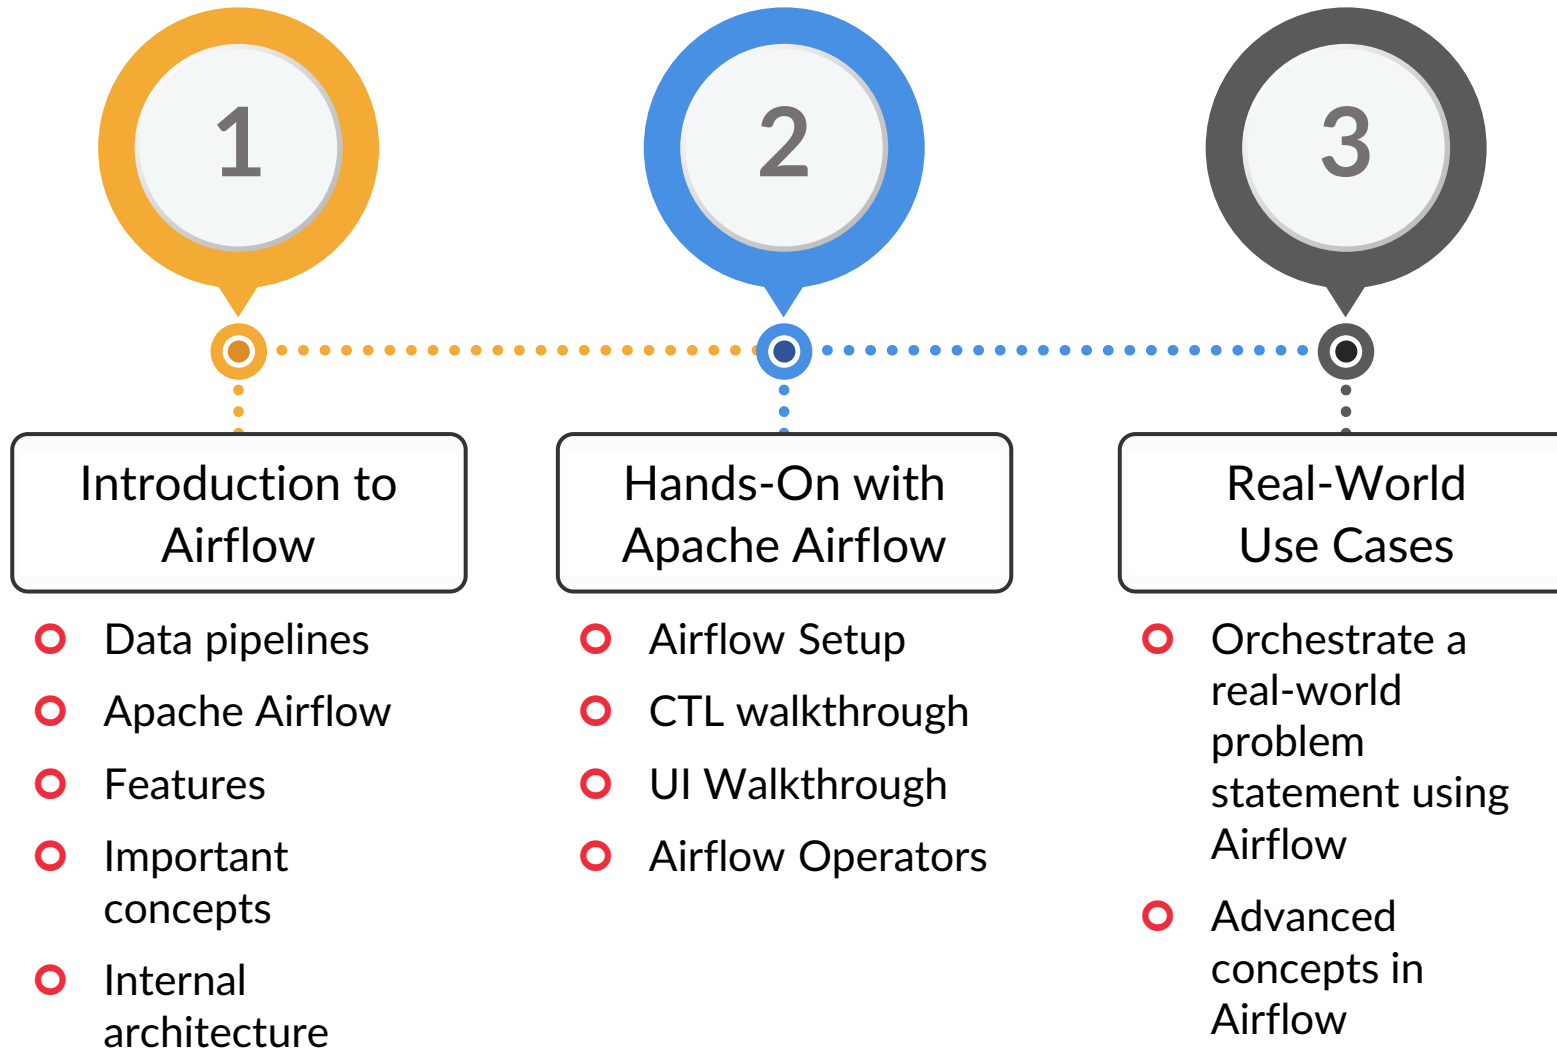


Introduction to Apache Airflow

Session 1

HEALTHCARE EVOLUTION



SESSION OVERVIEW

01

Data pipelines in large scale distributed environments

02

Solutions to data pipeline orchestration

03

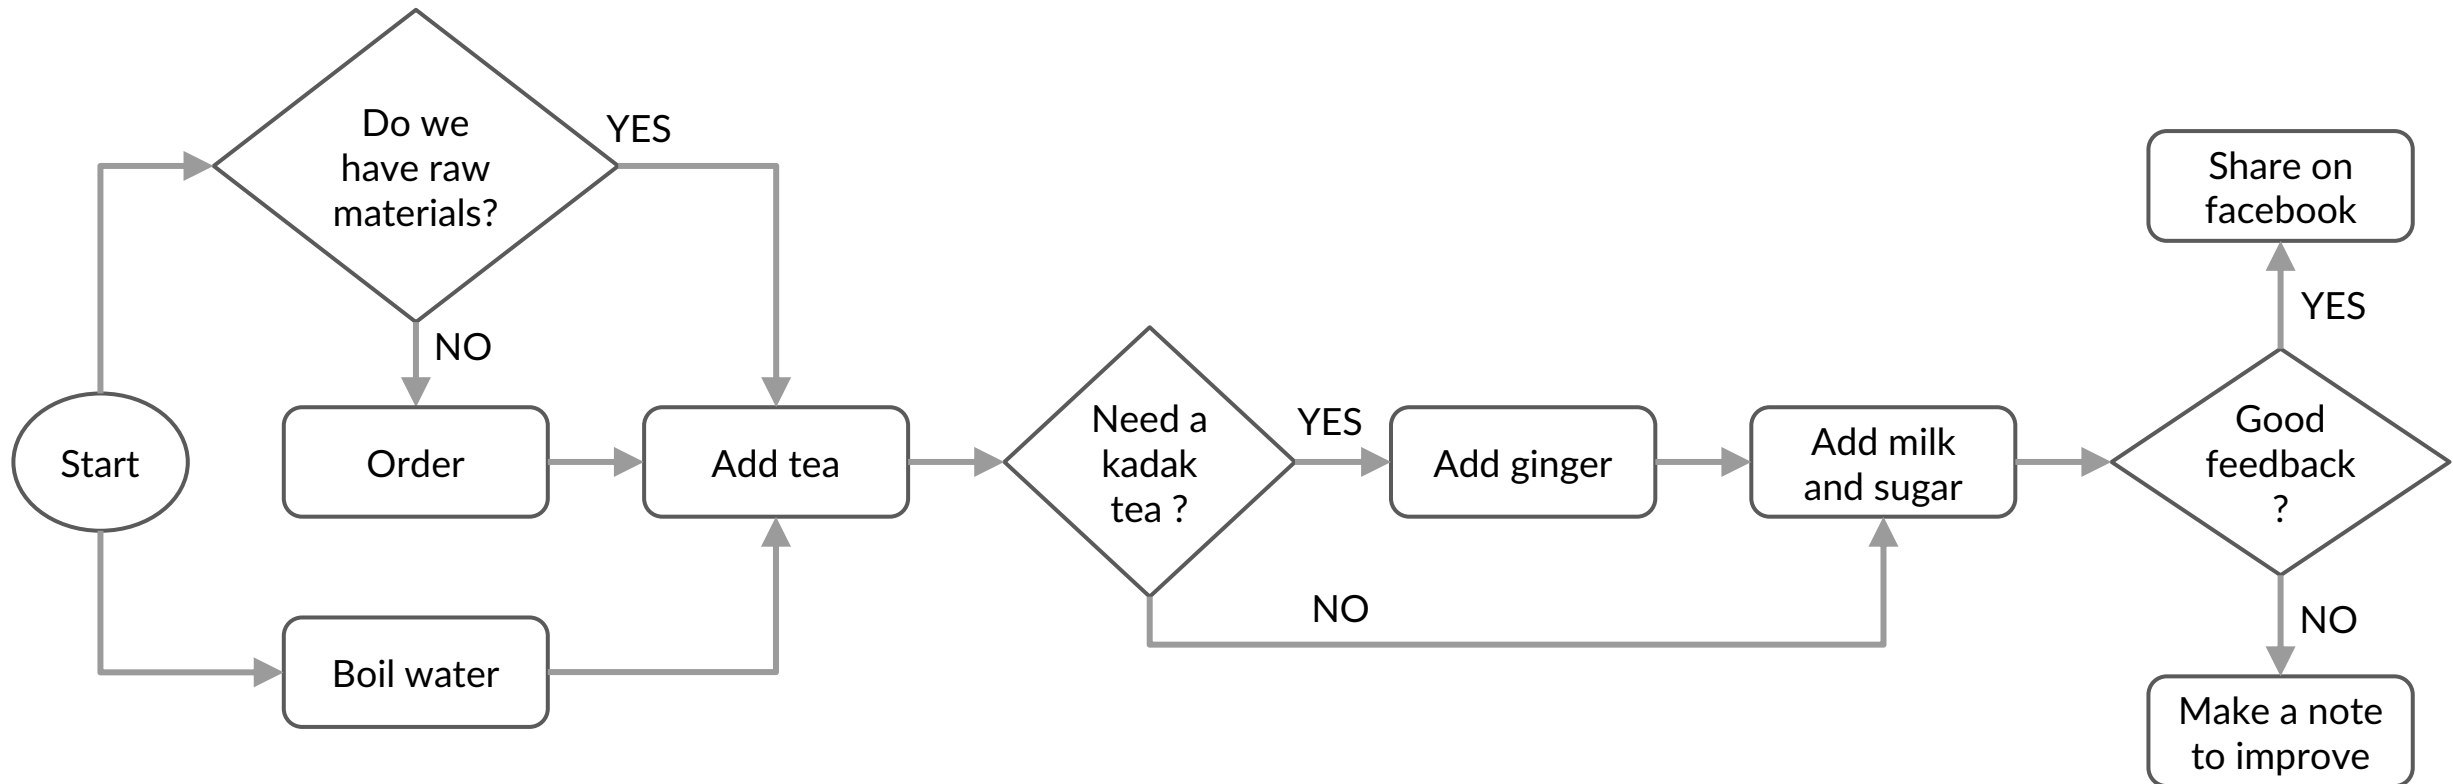
Apache Airflow and its features

04

The architecture of Airflow

FLOW CHART

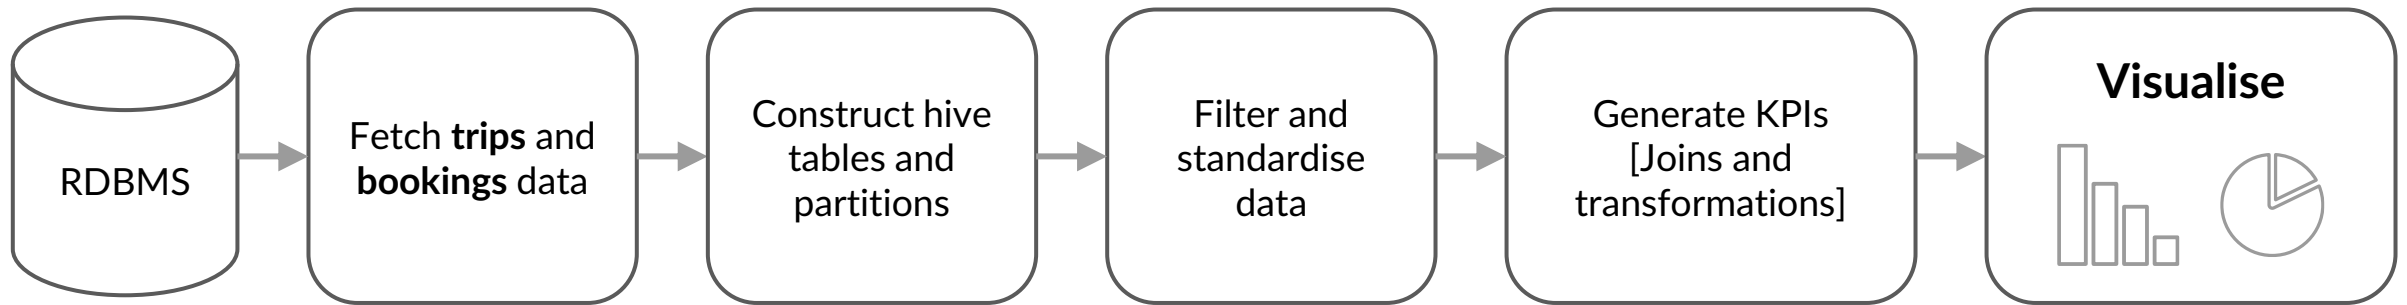
- Visual representation of the flow of actions taken to reach to an output or conclusion
- Easy to understand the logical steps happening in the activity



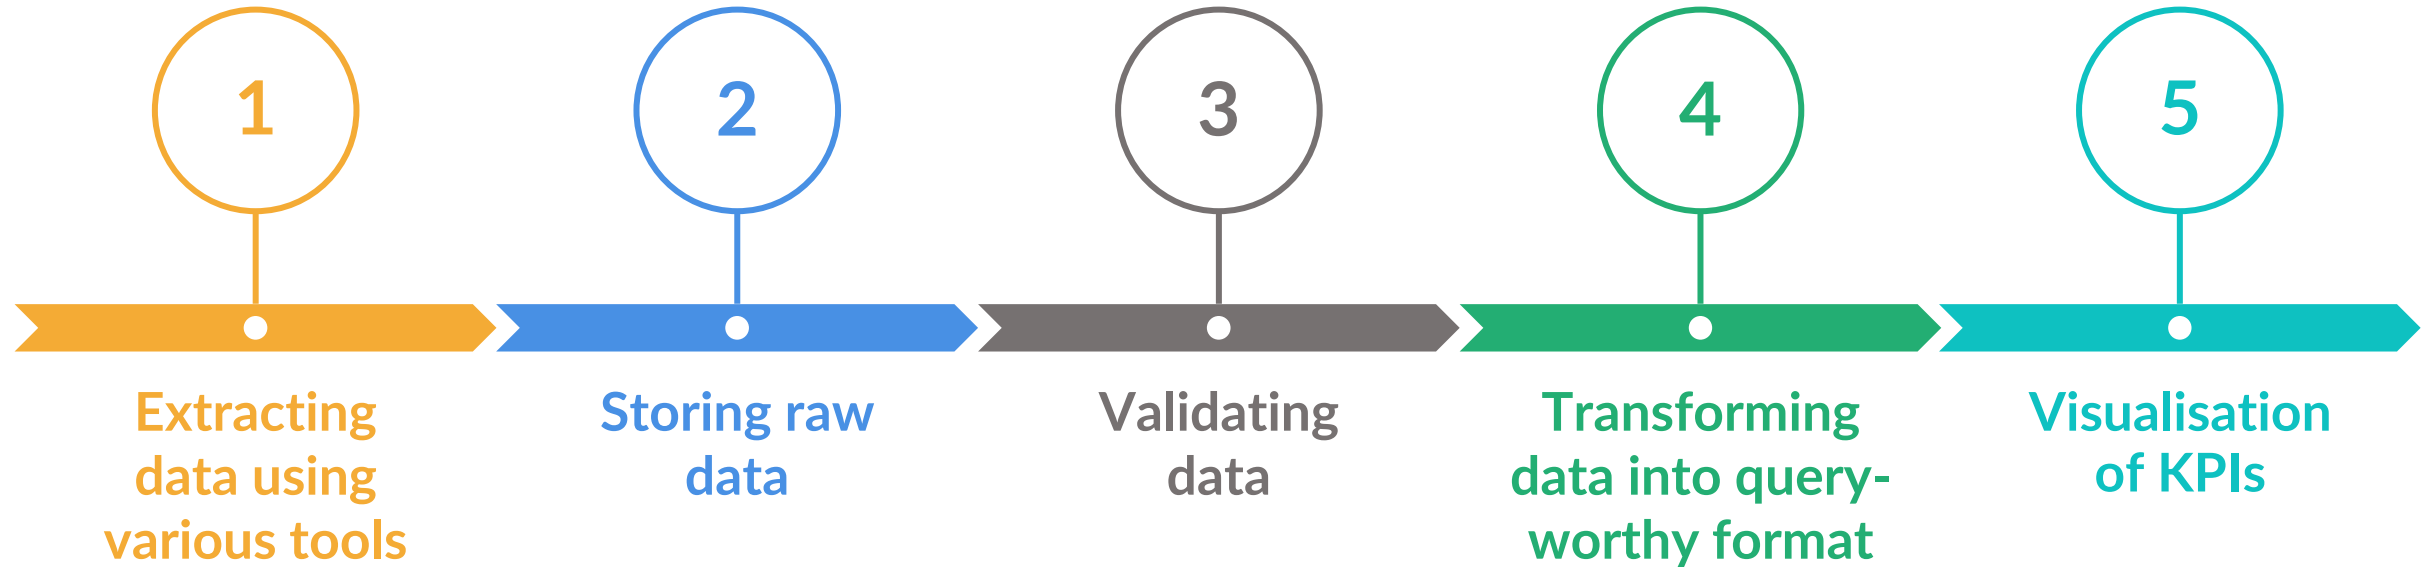
DATA PIPELINE

- Databases good for Online Transactional Processing (OLTP) are not best suited for Online Analytical Processing (OLAP)
- First step in cost effective and scalable solution for data processing is to bring data from OLTP databases to OLAP systems
- Data Pipeline
 - Process of moving data from one system to the other and possibly transforming and validating it in the process

UBER USE CASE



GENERAL STRUCTURE OF DATA PIPELINES



Orchestration of the
aforementioned
process

IMPORTANCE OF DATA PIPELINE AUTOMATION

- Data pipeline acts as the central nervous system for the **data-driven-decision-making**.
- **Notification** when processed data is available or if something fails
- **Reduced manual effort. More focus on business logic.**
- Keep track of the **performance** of different steps in the pipeline. This helps in **identification and resolution of bottlenecks**
- Automating data pipelines allows the company to collect, process and economically use data in **real-time**

POSSIBLE SOLUTIONS TO DATA PIPELINE AUTOMATION

- Manual Orchestration
 - No need to learn any data-pipeline orchestration tool
 - Impractical in a large scale industry environment
 - Against the principle of Do-Not-Repeat

POSSIBLE SOLUTIONS

- Cron Jobs
 - Very simple, comes pre-installed in Linux machines, easy to learn
 - Cumbersome when pipelines are longer
 - No intelligent handling of failed tasks or task retries
 - Lacks in-built alerting and monitoring
 - Not extensible
 - No visualization
 - Difficult to get info about the components of the pipeline or current status or previous runs

POSSIBLE SOLUTIONS

- Apache Oozie
 - One of the early and robust data orchestration tools with
 - ◆ scheduling capabilities
 - ◆ UI
 - ◆ parameterization (using variables)
 - ◆ supporting diverse operations (like Email, MapReduce, Hive etc.)
 - Steep learning curve
 - Uses XML for building pipelines, difficult to create dynamic DAGs
 - Difficult to build pipelines with a mediocre GUI
 - Community support is not great

WHY DO WE NEED A NEW TOOL?

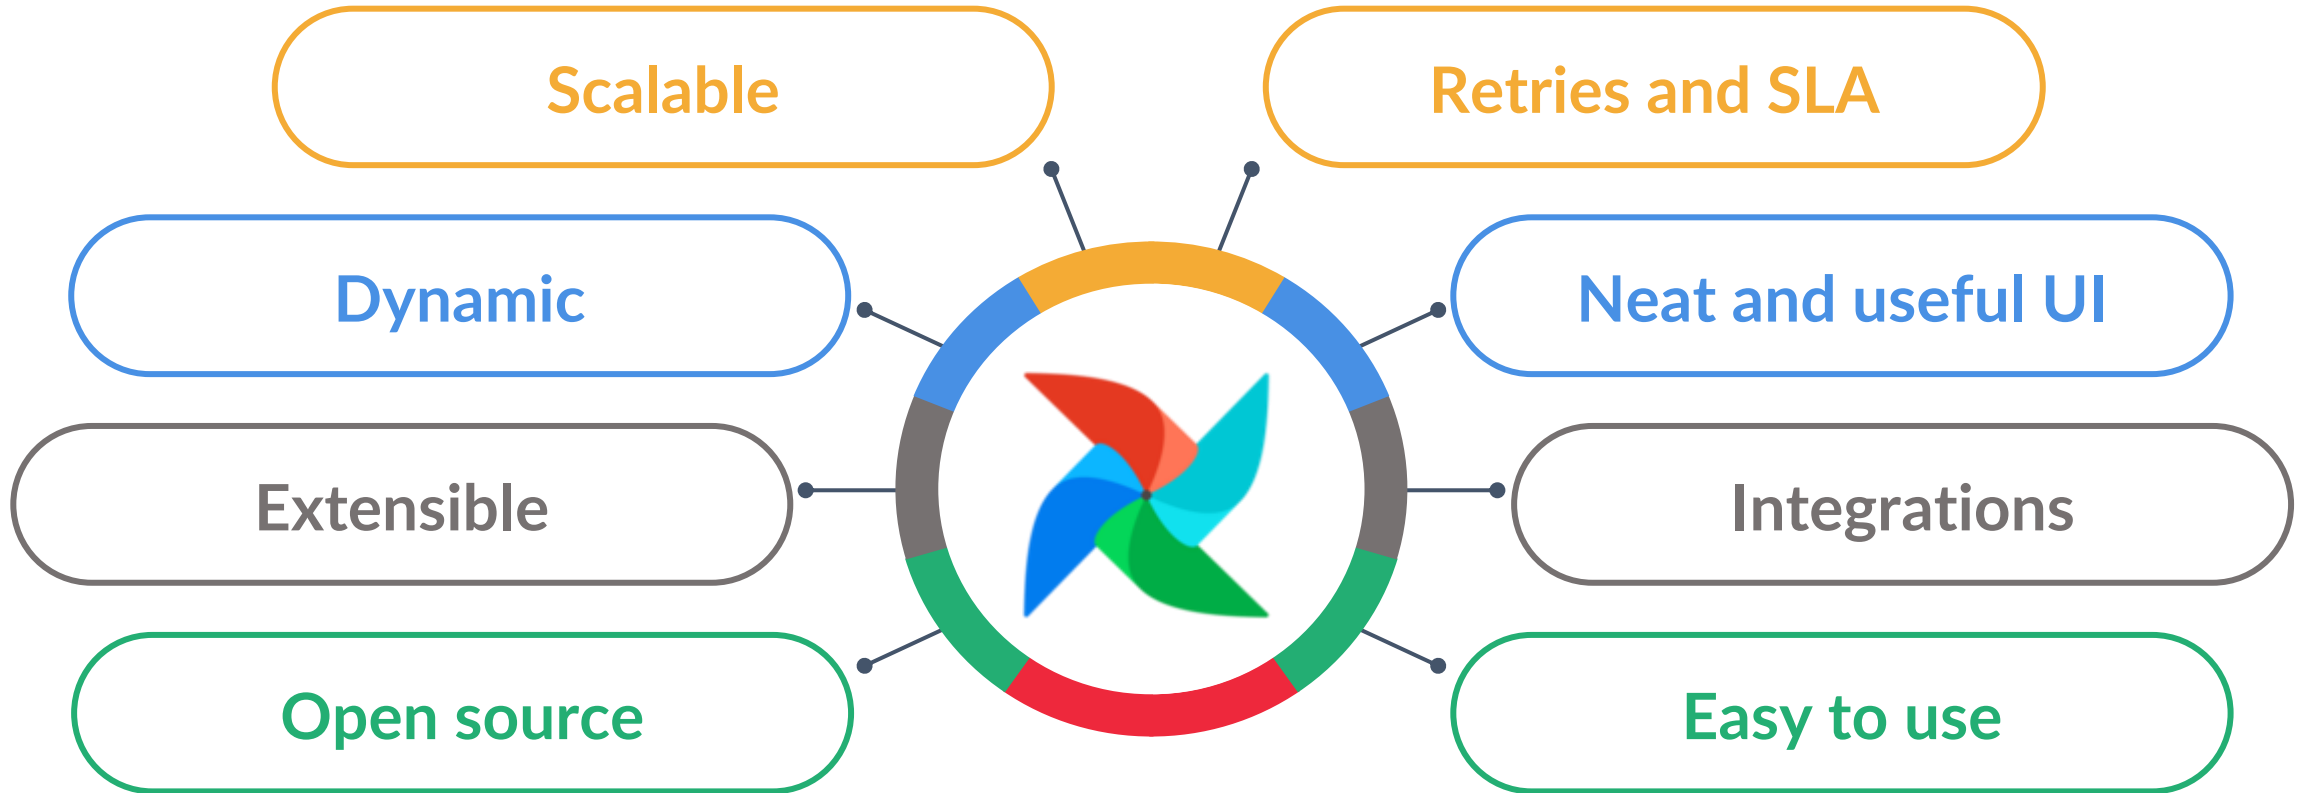
- Scalable to **orchestrate** and **schedule** large number of processes
- Efficient **maintenance** and **monitoring** of pipelines
- **Richer UI** to **visualise** the status and compare against historical runs
- **Retry** and **timeout** feature for failed tasks
- Easy **SLA** handling
- **Compatibility** with various tools in the data engineering domain
- **Strong community support**

INTRODUCTION TO APACHE AIRFLOW

- Started in October 2014 by Maxime Beauchemin at **Airbnb**
- Aims at meeting the needs of complex data processing pipelines
- Framework to programmatically **create, schedule and monitor** data pipelines
- Data pipelines in Airflow
 - Implemented in the form of **DAGs**
 - Created from the Python code
 - Can be generated dynamically



FEATURES OF APACHE AIRFLOW



DAG : DATA PIPELINES IN AIRFLOW

- DAGs or **Directed Acyclic Graphs** form the core structure around which Airflow structures its data pipelines.
- Properties of **DAGs** include:
 - Should be **directed** (i.e., the edges all have a direction indicating which task is dependent on which).
 - Must be **acyclic** (i.e., they can't contain cycles).
 - Are **graph** structures (i.e., a collection of vertices and edges).
- The whole DAG is defined in a Python program.

DAG ATTRIBUTES

- When creating a DAG we require :
 - ID
 - Description
 - Schedule (can be a cron expression or shorthand string notation like @once, @hourly etc.)
 - Start Date
 - Configs/default arguments
 - ◆ Owner
 - ◆ Retries


DAG DEFINITION

```
from datetime import datetime
from airflow import DAG

dag_default_configs = {
    'start_date': datetime(2016, 1, 1),
    'owner': 'airflow'
}

dag_object = DAG('my_dag',
                  default_args = dag_default_configs,
                  description='Sample DAG',
                  schedule_interval='0 12 * * *')
```

CRON EXPRESSION GUIDE



COMPONENTS OF A DAG

A DAG consists of **tasks** and **dependencies** between them.

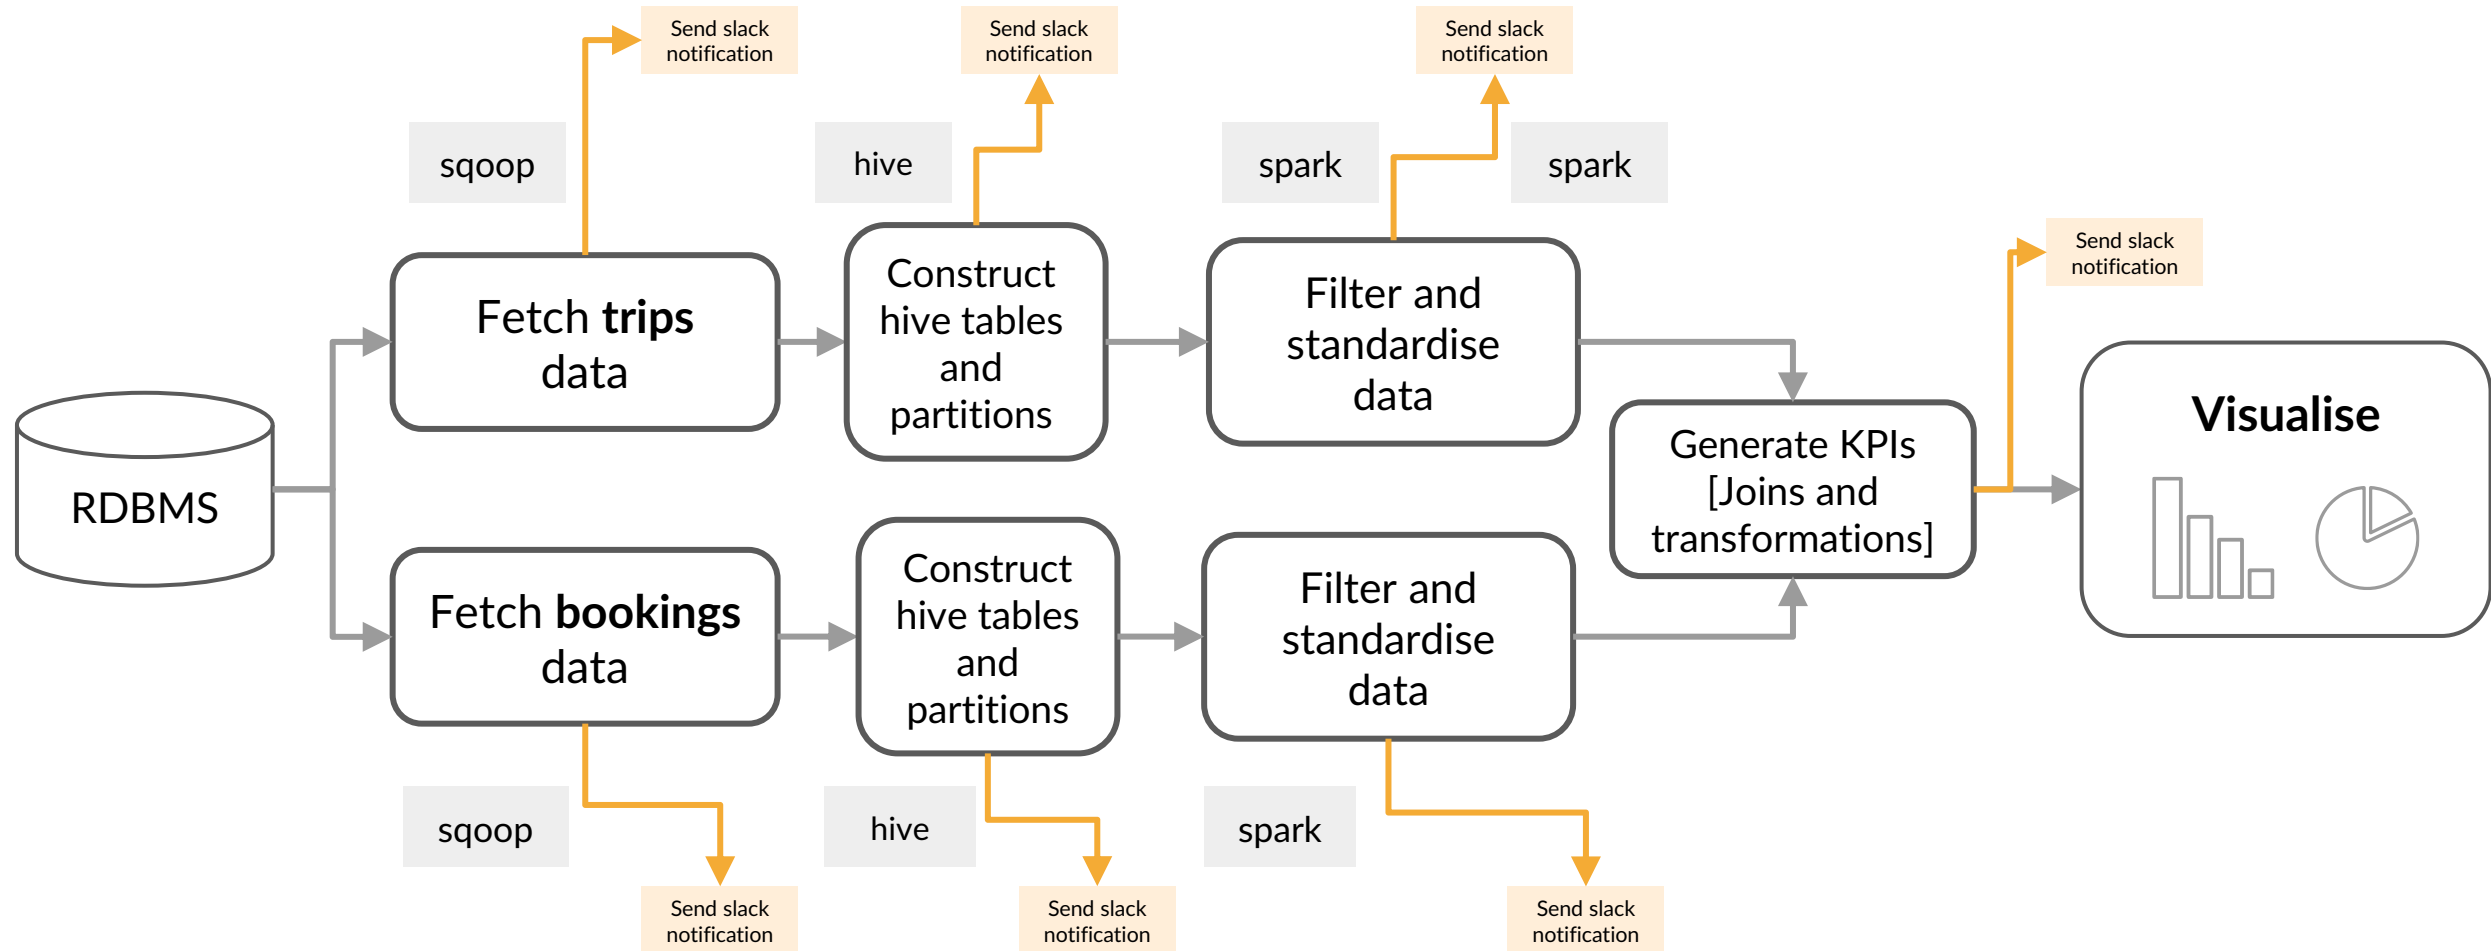
- Tasks :

- A Task defines a unit of work within a DAG
- It is represented as a node in the DAG graph
- Operators determine what actually gets done by a task.
- Each task is a python process running on the same or different machines (called worker)
- E.g: calling a python function, executing a shell script, running a hive query, running a spark job etc.

- Task dependencies :

- They define the order in which the tasks in a DAG are executed.

CONSTRUCTING DAG FOR UBER USE CASE



COMPONENTS OF AIRFLOW



Web server serves as the front-end



Scheduler orchestrates various DAGs and their tasks



Executor is the mechanism by which task instances get run

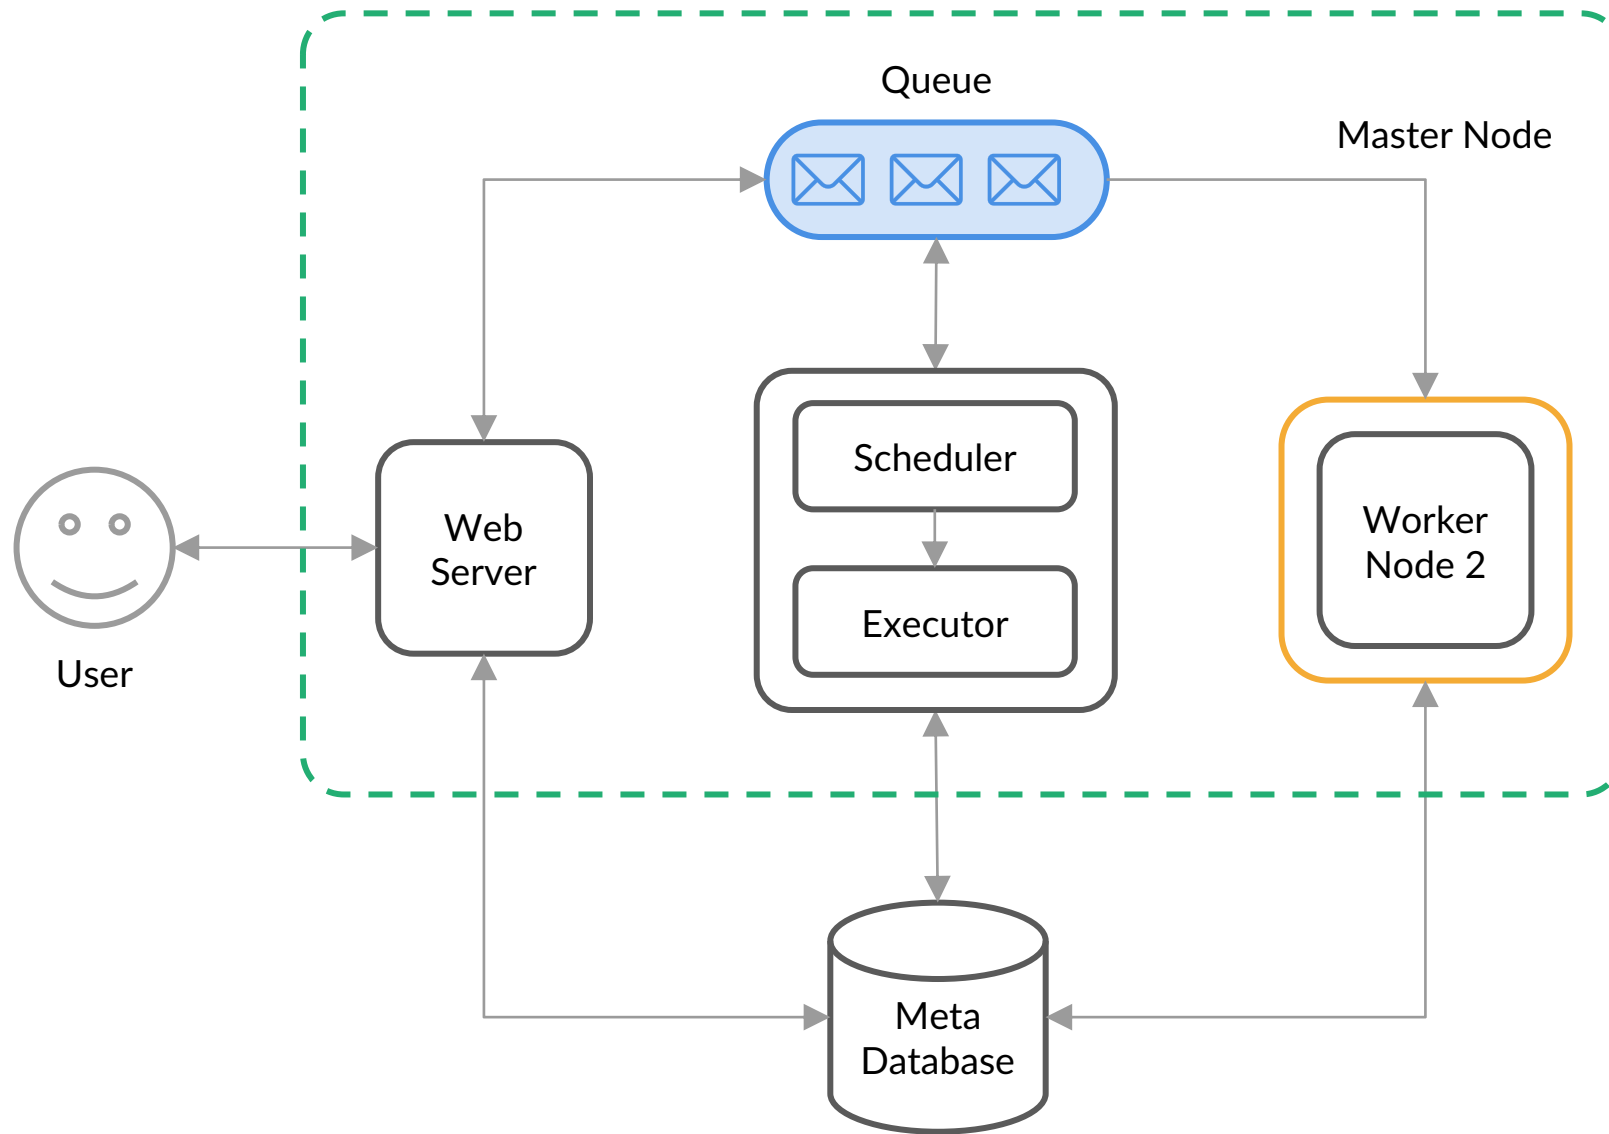


Workers perform the actual work



Metadata Database is used to keep track of task job statuses and other persistent(meta) information

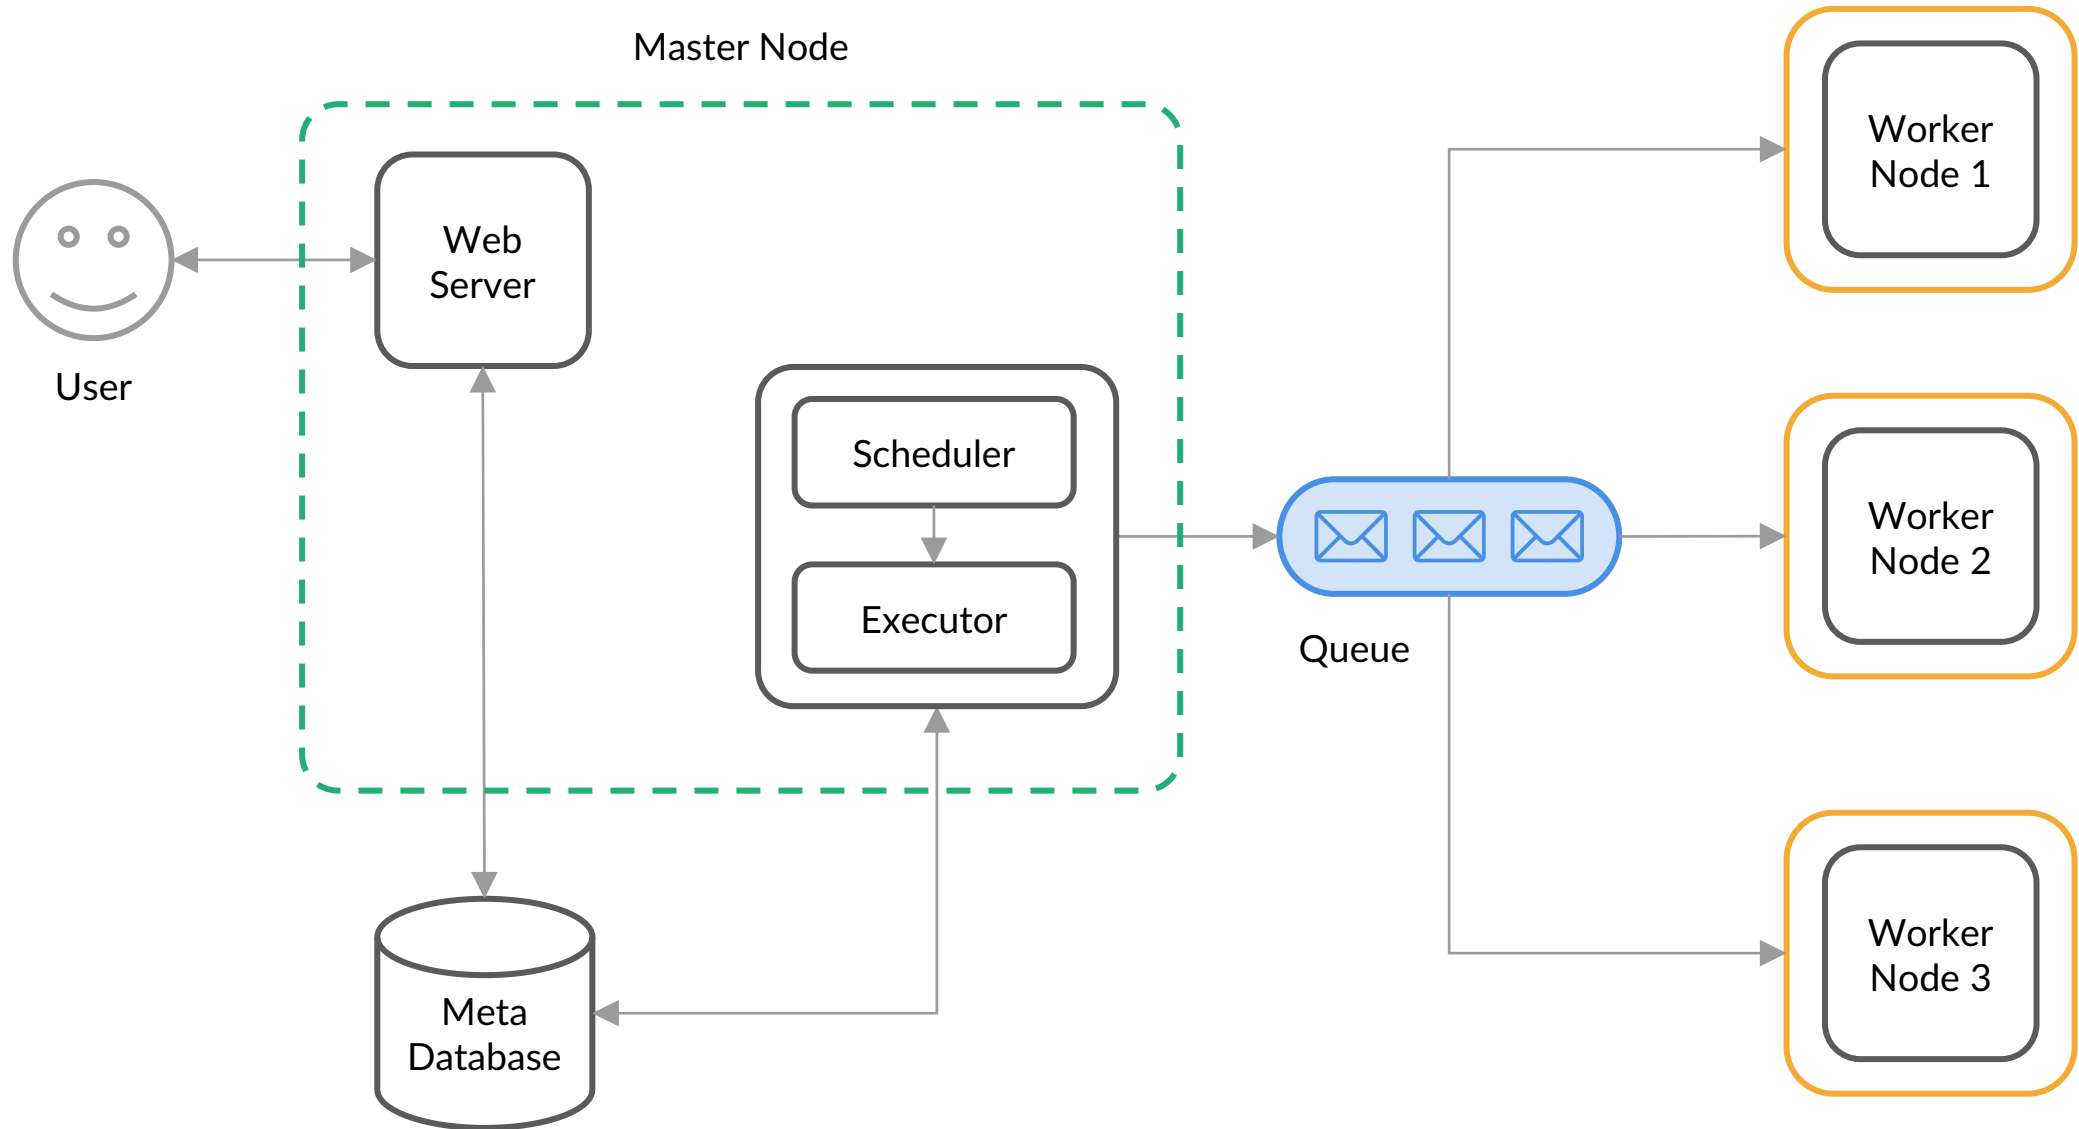
SINGLE NODE ARCHITECTURE



SCALING AIRFLOW: TYPES OF EXECUTORS

- Sequential executor:
 - It is the default executor and will only run one task instance at a time.
 - It is suitable for testing and debugging DAGs before they're implemented in an industry environment.
- Local executor:
 - It is like a sequential executor with unlimited parallelism.
 - It runs tasks by spawning processes in a controlled fashion in different modes.
- Celery executor:
 - Used in scalable environments
 - Needs RabbitMQ and Redis for configuration.
 - Each worker is in a different node, so it can be easily scaled by adding more nodes.
 - Recommended in production scenario

DISTRIBUTED ARCHITECTURE



SESSION SUMMARY

01

Data pipelines

02

Importance of pipeline orchestration

03

Possible solutions to data pipeline orchestration

04

Introduction to Airflow

05

Airflow Architecture

THANK YOU